

Graphs: Traversals and Shortest Path Algorithms (Chapter 9)

CSE 373
Data Structures and Algorithms

Today's Outline

- Announcements
 - Homework #6/7 ~~considered~~.
- **Graphs**
 - ~~Topological Sort~~
 - Shortest Paths Algorithms

Graph Traversals

- **Breadth-first search**
 - explore *all* adjacent nodes, then for each of *those* nodes explore *all* adjacent nodes
- **Depth-first search**
 - explore first child node, then its first child node, etc. until goal node is found or node has no children. Then backtrack, repeat with sibling.
- Both:
 - Work for arbitrary (directed or undirected) graphs
 - Must mark visited vertices so you do not go into an infinite loop!
- Either can be used to determine connectivity:
 - Is there a **path** between two given vertices?
 - Is the graph (weakly) connected?
- Which one:
 - Uses a queue? **BFS**
 - Uses a stack? **DFS**
 - Always finds the **shortest path** (for unweighted graphs)? **BFS**

The Shortest Path Problem

Given a graph G , edge costs c_{ij} , and vertices s and t in G , find the shortest path from s to t .

For a path $p = v_0 v_1 v_2 \dots v_k$

– *unweighted length* of path $p = k$ (a.k.a. *length*)

– *weighted length* of path $p = \sum_{i=0..k-1} c_{i,i+1}$ (a.k.a. *cost*)

Path length equals path cost when ?

Single Source Shortest Paths (SSSP)

Given a graph G , edge costs c_{ij} , and vertex s , find the shortest paths from s to all vertices in G .

– { Is this harder or easier than finding the shortest path from s to t ? } }

All Pairs Shortest Paths (APSP)

Given a graph G and edge costs $c_{i,j}$, find the shortest paths between all pairs of vertices in G .

- Is this harder or easier than SSSP?
- Could we use SSSP as a subroutine to solve this?

Variations of SSSP

- Weighted vs. unweighted
- Directed vs undirected
- Cyclic vs. acyclic
- Positive weights only vs. negative weights allowed
- Shortest path vs. longest path
- ...

Applications

- Network routing
- Driving directions
- Cheap flight tickets
- Critical paths in project management
(see textbook)
- ...

SSSP: Unweighted Version

Ideas?



10/EXP/DST/DEO/1/A

```
void Graph::unweighted (Vertex s){
    Queue q(NUM_VERTICES);
    Vertex v, w;
    for all vertices v do {v.dist = INFINITY;}
    s.dist = 0;
    q.enqueue(s);

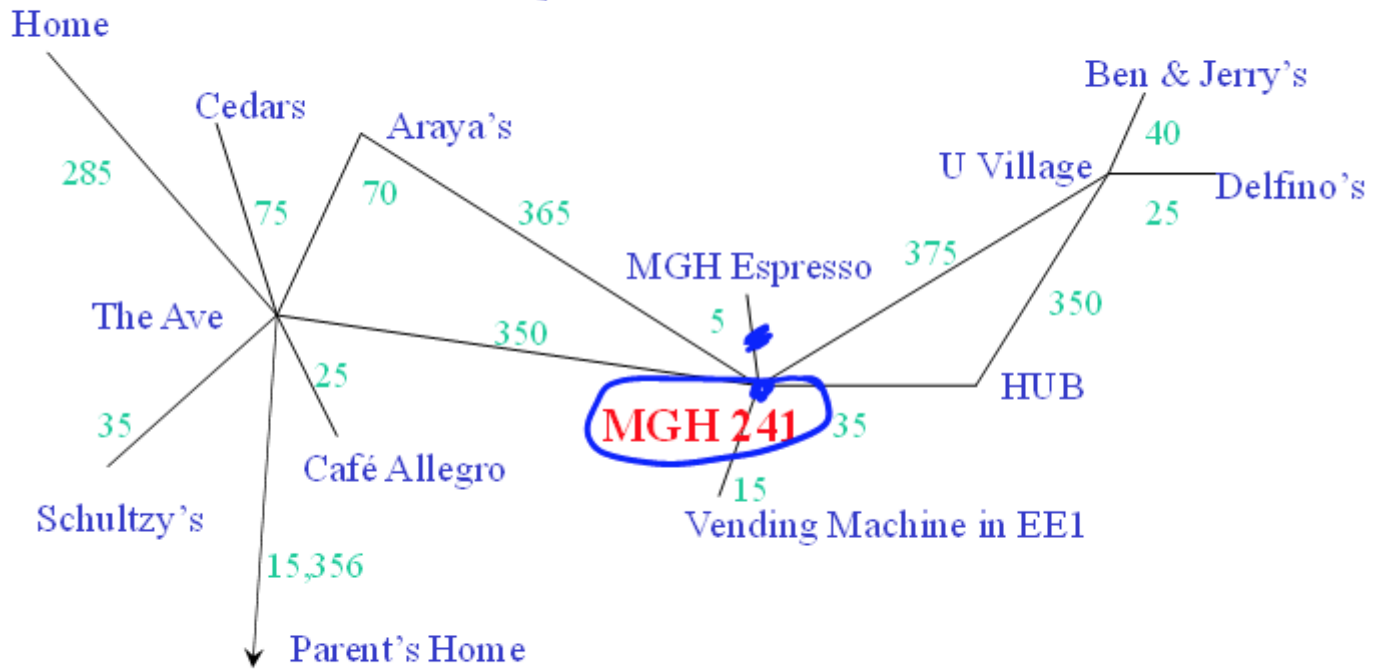
    while (!q.isEmpty()){
        v = q.dequeue();
        for each w adjacent to v
            if (w.dist == INFINITY){
                w.dist = v.dist + 1;
                w.path = v;
                q.enqueue(w);
            }
    }
}
```

each edge examined
at most once – if adjacency
lists are used

each vertex enqueued
at most once

total running time: $O(\quad)$

Weighted SSSP: The Quest For Food



Can we calculate shortest distance to all nodes from MGH 241?

Edsger Wybe Dijkstra
(1930-2002)




- Legendary figure in computer science; was a professor at University of Texas.
- Invented concepts of structured programming, synchronization, and "semaphores" for controlling computer processes.
- Supported teaching programming without computers (pencil and paper)
- 1972 Turing Award
- "In their capacity as a tool, computers will be but a ripple on the surface of our culture. In their capacity as intellectual challenge, they are without precedent in the cultural history of mankind."

Dijkstra's Algorithm for Single Source Shortest Path

- Similar to breadth-first search, but uses a heap instead of a queue:
 - Always select (expand) the vertex that has a lowest-cost path ~~to~~^{from} the start vertex
- Correctly handles the case where the lowest-cost (shortest) path to a vertex is not the one with fewest edges

Dijkstra's Algorithm: Idea




Adapt BFS to handle weighted graphs

Two kinds of vertices:

- Finished or **known** vertices
 - Shortest distance has been computed
- **Unknown** vertices
 - Have tentative distance

Dijkstra's Algorithm: Idea



At each step:

- 1) Pick closest **unknown** vertex
- 2) Add it to **known** vertices
- 3) Update distances

Dijkstra's Algorithm: Pseudocode

Initialize the cost of each node to ∞

Initialize the cost of the source to 0

While there are **unknown** nodes left in the graph

{ Select an **unknown** node b with the lowest cost

Mark b as **known**

For each node a adjacent to b

a 's cost = \min (a 's old cost, b 's cost + cost of (b, a))

17/EXP/ALG/DEO/1/A

```
void Graph::dijkstra(Vertex s){
    Vertex v,w;


    Initialize s.dist = 0 and set dist of all other
    vertices to infinity

    while (there exist unknown vertices, find the
    one b with the smallest distance)
        b.known = true;

        for each a adjacent to b
            if (!a.known)
                if (b.dist + Cost_ba < a.dist){
                    decrease(a.dist to= b.dist + Cost_ba);
                    a.path = b;
                }
        }
    }
```

Important Features

- Once a vertex is made **known**, the cost of the shortest path to that node is known
- While a vertex is still not **known**, another shorter path to it might still be found
- The shortest path itself can found by following the backward pointers stored in **node.path**



V	Known	Dist	path
v0	T	0	
v1	F	6	v3
v2	F	2	v0
v3	T	1	v0
v4	F	2	v3
v5	F	7	v3
v6	F	6	v3

Order Known
v0 v3

Dijkstra's Alg: Implementation

Initialize the cost of each node to ∞

Initialize the cost of the source to 0

While there are unknown nodes left in the graph

 Select the unknown node b with the lowest cost

 Mark b as known

 For each node a adjacent to b

a 's cost = $\min(a$'s old cost, b 's cost + cost of (b, a))

Running time?

Dijkstra's Algorithm: a Greedy Algorithm


Greedy algorithms always make choices that *currently* seem the best

- Short-sighted – no consideration of long-term or global issues
- Locally optimal - does not always mean globally optimal!!

Dijkstra's Algorithm: Summary

- Classic algorithm for solving SSSP in weighted graphs *without negative weights*
- A *greedy* algorithm (irrevocably makes decisions without considering future consequences)
- Intuition for correctness:
 - shortest path from source vertex to itself is 0
 - cost of going to adjacent nodes is at most edge weights
 - cheapest of these must be shortest path to that node
 - update paths for new node and continue picking cheapest path

Correctness: The Cloud Proof



How does Dijkstra's decide which vertex to add to the Known set next?

- If path to **V** is shortest, path to **W** must be *at least as long*
(or else we would have picked **W** as the next vertex)
- So the path through **W** to **V** cannot be any shorter!

Correctness: Inside the Cloud

Prove by induction on # of nodes in the cloud:

Initial cloud is just the source with shortest path 0

Assume: Everything inside the cloud has the correct shortest path

Inductive step: Only when we prove the shortest path to some node v (which is not in the cloud) is correct, we add it to the cloud

When does Dijkstra's algorithm not work?

Dijkstra's vs BFS

At each step:

- 1) Pick **closest unknown** vertex
- 2) Add it to finished vertices
- 3) Update distances

Dijkstra's Algorithm

At each step:

- 1) Pick vertex from **queue**
- 2) Add it to visited vertices
- 3) Update queue with neighbors

Breadth-first Search

26/STR/ALG/DEF/1/A

The Trouble with Negative Weight Cycles

What's the shortest path from A to E?

Problem?