

AVL-Trees



Balanced binary tree

- The disadvantage of a binary search tree is that its height can be as large as $N-1$
- This means that the time needed to perform insertion and deletion and many other operations can be $O(N)$ in the worst case
- We want a tree with small height
- A binary tree with N node has height **at least** $\Theta(\log N)$
- Thus, our goal is to keep the height of a binary search tree $O(\log N)$
- Such trees are called **balanced** binary search trees. Examples are AVL tree, red-black tree.



AVL tree

Height of a node


- ☞ The height of a leaf is 1. The height of a null pointer is zero.
- ☞ The height of an internal node is the maximum height of its children plus 1

Note that this definition of height is different from the one we defined previously (we defined the height of a leaf as zero previously).



AVL tree

- An AVL tree is a binary search tree in which
 - for *every* node in the tree, the height of the left and right subtrees differ by **at most 1**.



AVL property
violated here




Figure 4.32 Two binary search trees. Only the left tree is AVL.



AVL tree

- Let x be the root of an AVL tree of height h
- Let N_h denote the minimum number of nodes in an AVL tree of height h
- Clearly, $N_i \geq N_{i-1}$ by definition
- We have

$$\begin{aligned}
 N_h &\geq N_{h-1} + N_{h-2} + 1 \\
 &\geq 2N_{h-2} + 1 \\
 &> 2N_{h-2}
 \end{aligned}$$
- By repeated substitution, we obtain the general form



$$N_h > 2^i N_{h-2}$$
- The boundary conditions are: $N_1=1$ and $N_2=2$. This implies that $h = O(\log N_h)$.
- Thus, many operations (searching, insertion, deletion) on an AVL tree will take **$O(\log N)$ time**.

Rotations



- When the tree structure changes (e.g., insertion or deletion), we need to transform the tree to restore the AVL tree property.
- This is done using **single rotations** or **double rotations**.

e.g. Single Rotation





Rotations

- ☞ Since an insertion/deletion involves adding/deleting a single node, this can only increase/decrease the height of some subtree by 1
- ☞ Thus, if the AVL tree property is violated at a node x , it means that the heights of $\text{left}(x)$ and $\text{right}(x)$ **differ by exactly 2**.
- ☞ Rotations will be applied to x to restore the AVL tree property.



Insertion

- First, insert the new key as a new leaf just as in ordinary binary search tree
- Then trace the path **from the new leaf towards the root**. For each node x encountered, check if heights of $\text{left}(x)$ and $\text{right}(x)$ differ by at most 1.
- If yes, proceed to $\text{parent}(x)$. If not, restructure by doing **either a single rotation or a double rotation** [next slide].
- For insertion, once we perform a rotation at a node x , we won't need to perform any rotation at any ancestor of x .



Insertion

- ☞ Let x be the node at which $\text{left}(x)$ and $\text{right}(x)$ differ by more than 1
- ☞ Assume that the height of x is $h+3$
- ☞ There are 4 cases
 - Height of $\text{left}(x)$ is $h+2$ (i.e. height of $\text{right}(x)$ is h)
 - ▢ Height of $\text{left}(\text{left}(x))$ is $h+1 \Rightarrow$ single rotate with left child
 - ▢ Height of $\text{right}(\text{left}(x))$ is $h+1 \Rightarrow$ double rotate with left child
 - Height of $\text{right}(x)$ is $h+2$ (i.e. height of $\text{left}(x)$ is h)
 - ▢ Height of $\text{right}(\text{right}(x))$ is $h+1 \Rightarrow$ single rotate with right child
 - ▢ Height of $\text{left}(\text{right}(x))$ is $h+1 \Rightarrow$ double rotate with right child

Note: Our test conditions for the 4 cases are different from the code shown in the textbook. These conditions allow a uniform treatment between insertion and deletion.




Single rotation

The new key is inserted in the subtree A.

The AVL-property is violated at x

- height of left(x) is $h+2$
- height of right(x) is h .




Rotate with left child



Single rotation



The new key is inserted in the subtree C.
The AVL-property is violated at x.




Rotate with right child

Single rotation takes $O(1)$ time.
Insertion takes $O(\log N)$ time.

AVL Tree



Insert 0.8



After rotation




Double rotation

The new key is inserted in the subtree B1 or B2.

The AVL-property is violated at x.

x-y-z forms a zig-zag shape




also called left-right rotate



Double rotation

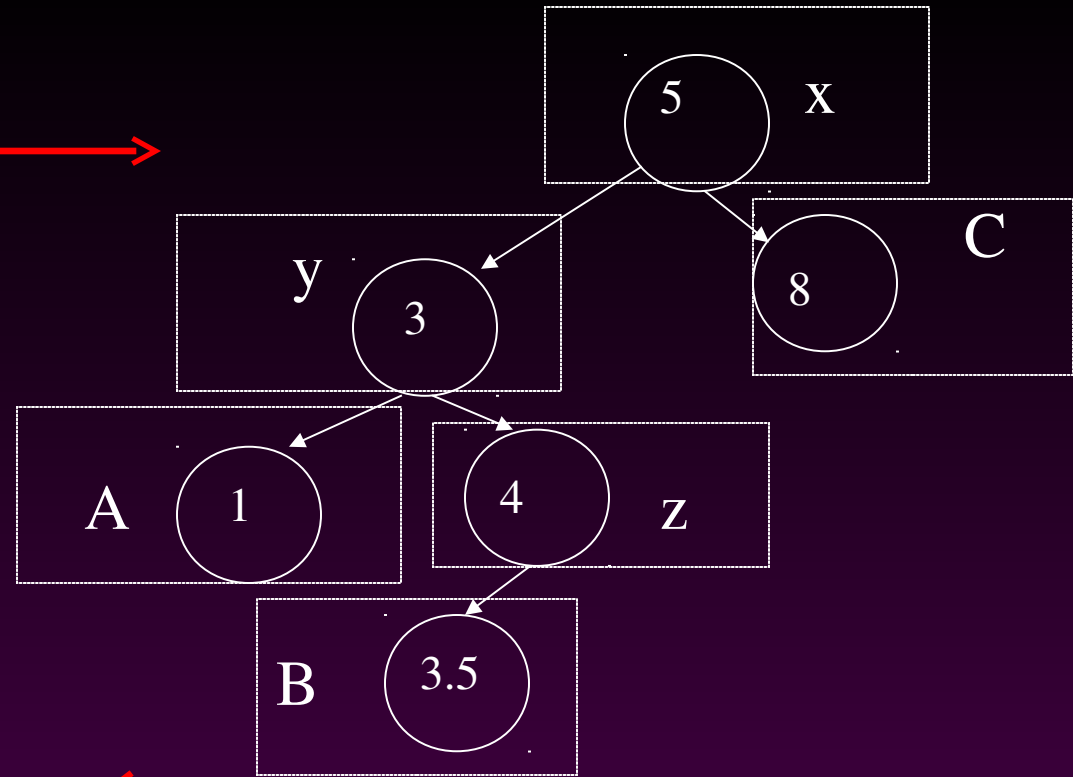

The new key is inserted in the subtree B1 or B2.
The AVL-property is violated at x.




Double rotate with right child

also called right-left rotate

AVL Tree



Insert 3.5



After Rotation



An Extended Example

Insert 3,2,1,4,5,6,7, 16,15,14




Fig 1




Fig 2




Fig 3




Fig 4




Fig 5




Fig 6

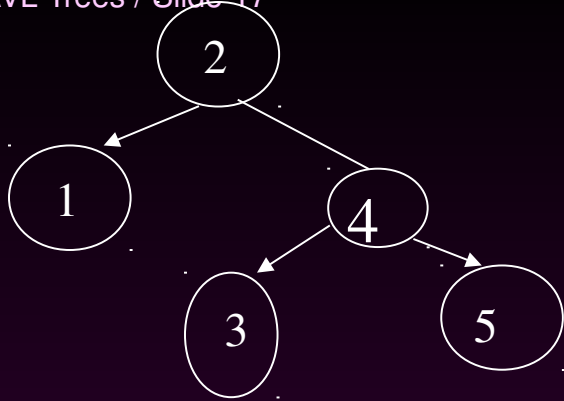


Fig 7

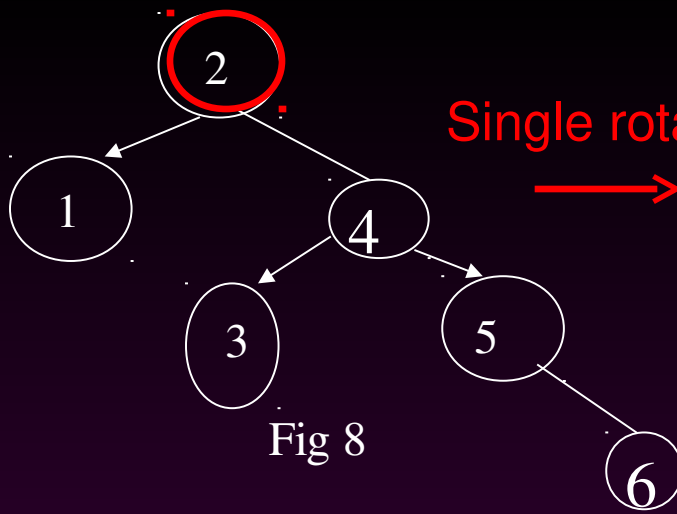


Fig 8

Single rotation

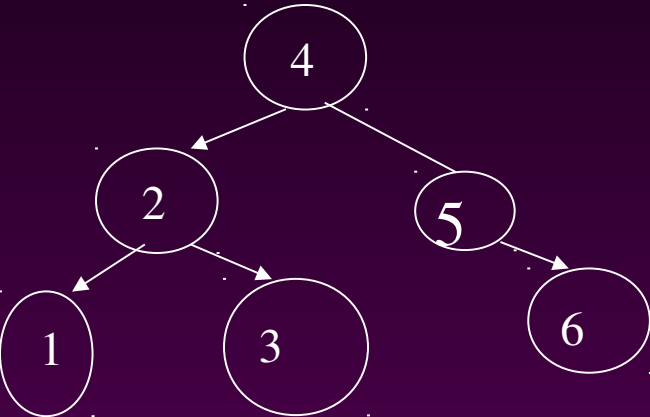


Fig 9

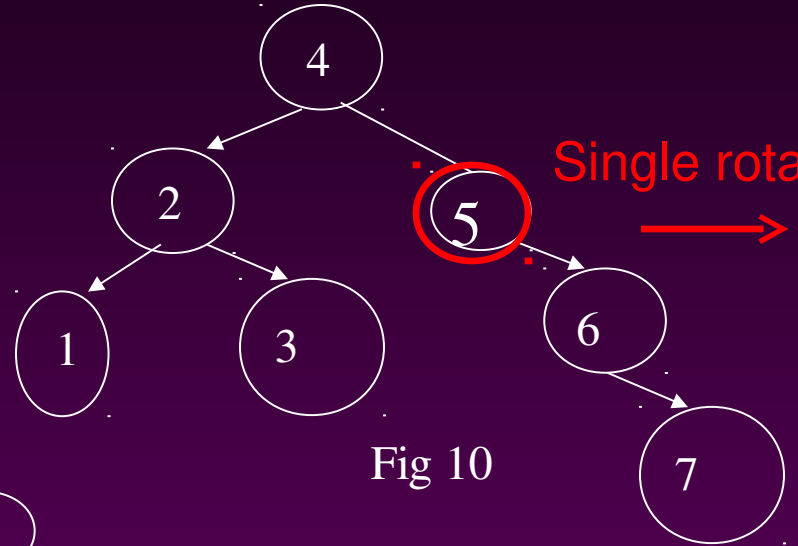


Fig 10

Single rotation

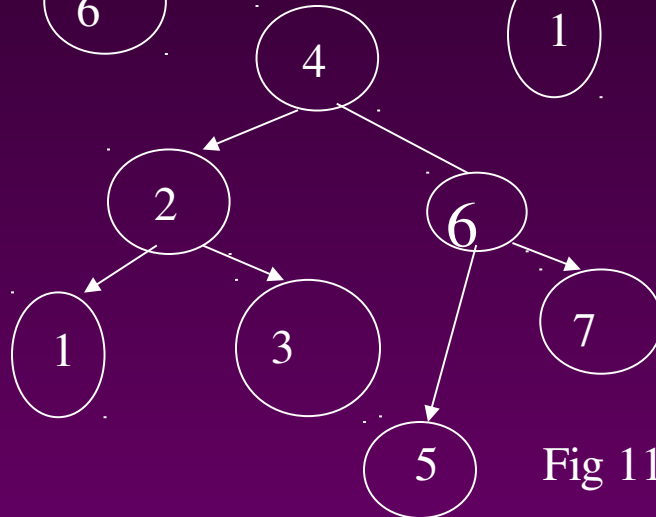


Fig 11




Fig 12




Fig 13

Double rotation




Fig 14

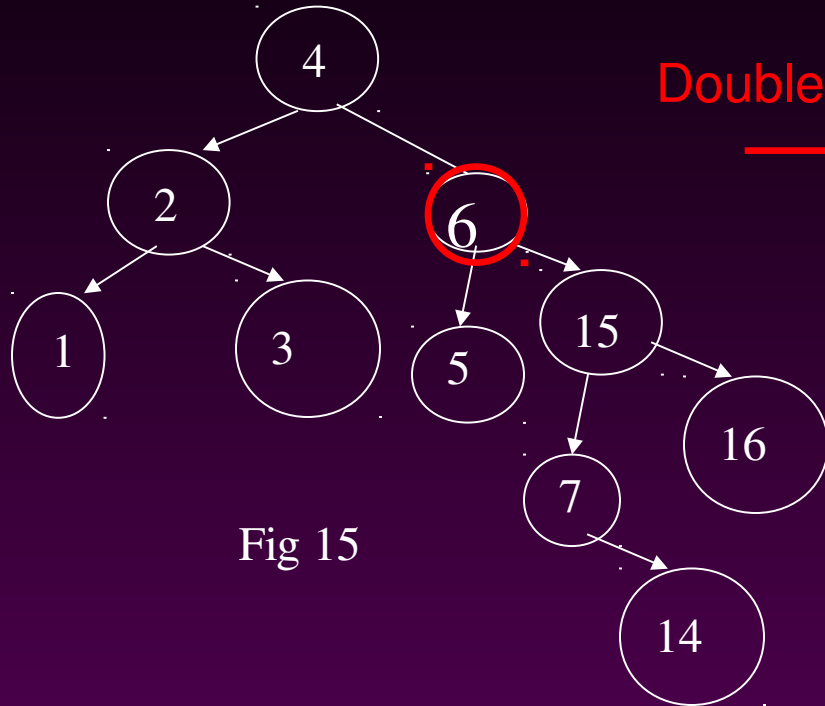


Fig 15

Double rotation

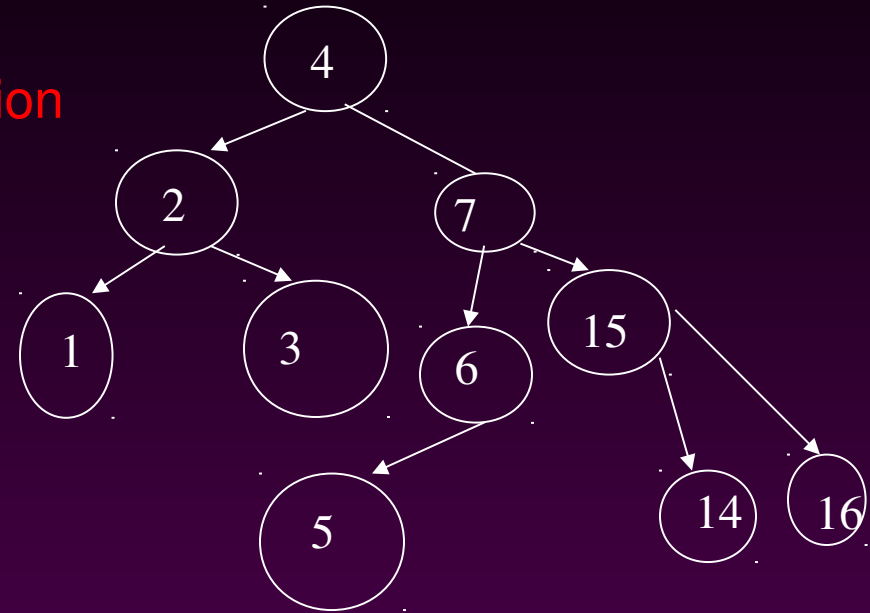


Fig 16



Deletion

- Delete a node x as in ordinary binary search tree. Note that the last node deleted is a leaf.
- Then trace the path from **the new leaf towards the root**.
- For each node x encountered, check if heights of $\text{left}(x)$ and $\text{right}(x)$ differ by at most 1. If yes, proceed to $\text{parent}(x)$. If not, perform an appropriate rotation at x . There are 4 cases as in the case of insertion.
- For deletion, after we perform a rotation at x , we may have to perform a rotation at some ancestor of x . Thus, we must **continue to trace the path until we reach the root**.

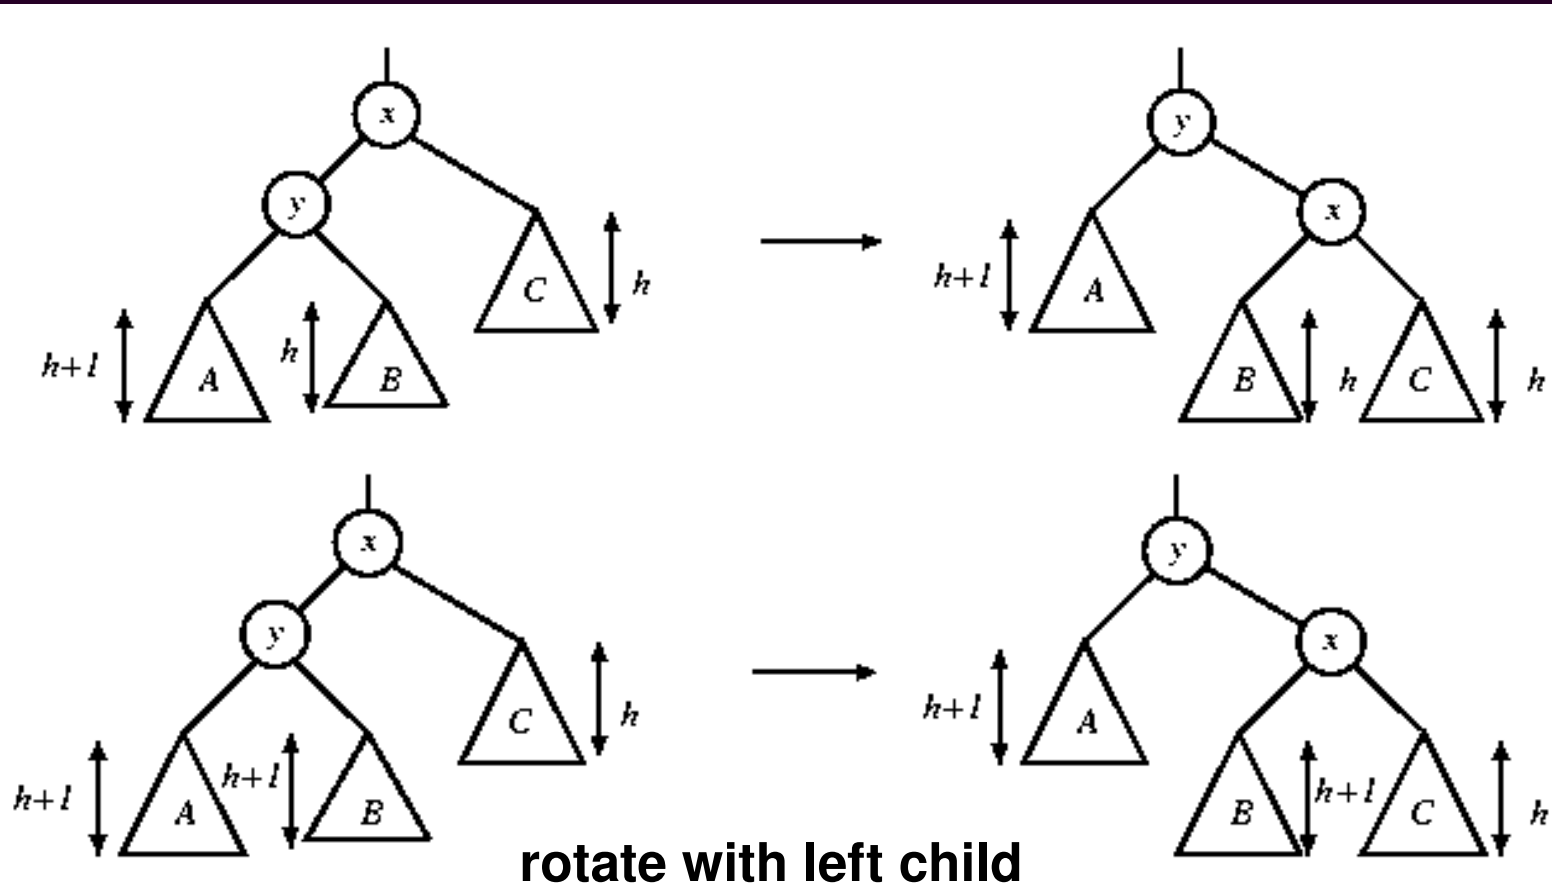


Deletion

- ☞ On closer examination: the single rotations for deletion can be divided into 4 cases (instead of 2 cases)
 - Two cases for rotate with left child
 - Two cases for rotate with right child

Single rotations in deletion



In both figures, a node is deleted in subtree C, causing the height to drop to h . The height of y is $h+2$. When the height of subtree A is $h+1$, the height of B can be h or $h+1$. Fortunately, the same single rotation can correct both cases.





Single rotations in deletion

In both figures, a node is deleted in subtree A, causing the height to drop to h . The height of y is $h+2$. When the height of subtree C is $h+1$, the height of B can be h or $h+1$. A single rotation can correct both cases.



Rotations in deletion

- There are 4 cases for single rotations, but we do not need to distinguish among them.
- There are exactly two cases for double rotations (as in the case of insertion)
- Therefore, we can reuse exactly the same procedure for insertion to determine which rotation to perform